

Saarthi 2.0

Detailed Design Document

Module flows, data model & table usage

Repository: TathyumSolutions/SAARTHI2.0

Source: docs/SAARTHI_DESIGN_DOCUMENT.md

Date: October 2026

Saarthi 2.0 — Detailed Design Document

Scope. This covers the whole Saarthi codebase as it stands on this branch. It focuses on:

1. **Runtime flows.** Which modules run, and in what order, for each user action. That means the chat question, plus every "setup" action such as **Save, Process, Upload, Like/Dislike** and **Grant access**.
2. **Data ownership.** Which database tables (and non-database stores) each flow **creates, reads, updates or deletes**, and how a later flow uses what an earlier flow wrote.

`docs/CODE_FLOW_QUESTION_ANSWERING.md` is a lighter walkthrough of the chat path only. This document is a superset of it and adds the data-model view.

Table of contents


1. System overview
2. Runtime components & deployment
3. Code layout (module map)
4. Persistence model
5. Tenancy & visibility rule ("own + granted")
6. MetaMind — the per-user router config
7. Setup flows (what populates the tables)
8. Query flow — what happens when a question is asked
9. Track internals
10. Self-learning (feedback & query reuse)
11. Observability: Query log, LLM call log, Audit log
12. Model selection & LLM backends
13. Data warehouse builder (side feature)
14. Table × flow CRUD matrix
15. Configuration reference
16. Observed gaps & risks in the current code

1. System overview

Saarthi is a multi-tenant "ask your company data" assistant. A user types a natural-language question. An LLM **router** then decides which of the user's connected data sources can answer it, dispatches the question to one or more **tracks**, merges the results, and returns an answer with an optional table and chart.

Track	Data it answers from	Engine
DB	External PostgreSQL databases the user registered (e.g. an SAP replica)	8-agent LangGraph pipeline → SQL
SPREADSHEET	Uploaded Excel/CSV files, stored as Parquet (never loaded into Postgres)	LLM writes a JSON query plan → pandas executes it
FILES	Uploaded documents (PDF/DOCX/MD/TXT ...)	RAG over Qdrant
API	Registered external REST endpoints	LLM tool-calling → HTTP request
GENERAL	No company data (world knowledge, greetings, dates)	Plain LLM call

A data source is only answerable after it goes through **setup**: register/upload → **Process** (introspect, summarise, index) → optionally **share** with colleagues through Resource Mapping. Setup writes the metadata that the router reads at question time. That metadata is called *MetaMind*.



2. Runtime components & deployment

`docker-compose.yml` runs these services:

Service	Image	Used for
web	Flask app (<code>run.py</code> → <code>app.create_app</code>)	All HTTP routes, HTML templates, every service in the same process
db	<code>postgres:15</code>	The app's own 3 logical databases (core / resources / workspace)
redis	<code>redis:7-alpine</code>	Flask-Limiter rate-limit storage (<code>RATELIMIT_STORAGE_URI</code> = <code>REDIS_URL</code>)
qdrant	<code>qdrant/qdrant</code>	Vector store for document chunks (collection <code>saarthi_unstructured</code>)
ollama (×2)	<code>ollama/ollama</code>	Local LLMs (<code>llama3</code> etc.) for the DB agents and for "OSS" model choices

Volumes: `postgres_data`, `redis_data`, `ollama_data`, `uploads_data` (uploaded files, Parquet spreadsheets, manifest), `nltk_data`, `hf_cache` (HuggingFace embedding model).

External, user-supplied systems: the **customer PostgreSQL databases** registered as connections, and the **REST APIs** registered as API connectors. Saarthi never writes business data into its own databases. It reads from these external systems at question time.

App bootstrap (`app/__init__.py:create_app`)

```

create_app(config_name)                                # run.py reads FLASK_ENV
├─ load config/config.py (derives 3 DB URLs from DATABASE_URL)
├─ db_bootstrap.bootstrap_databases(SQLALCHEMY_BINDS)  # CREATE DATABASE saarthi_core_db /
  _resources_db / _workspace_db if missing
├─ init SQLAlchemy, Migrate, JWT, Limiter(Redis), Swagger
├─ register 23 blueprints (pages + /api/*)
├─ import app.models (registers all ORM classes)
└─ app_context:
    ├── configure_mappers(); db.create_all()           # creates missing tables in each bind
    ├── db_bootstrap.sync_missing_columns(db)          # additive ALTER TABLE ADD COLUMN for
    new model fields
    └─ db_bootstrap.drop_removed_tables(db, ['router_configs']) # legacy cached-router table is
    dropped

```

The schema is defined by the ORM models; there is no hand-written DDL. Alembic is wired up (`migrations/`), but in practice `create_all` + `sync_missing_columns` keeps the schema in step.


3. Code layout (module map)

run.py	entry point (Flask dev server / gunicorn target)
config/config.py	env-driven config; 3 SQLAlchemy binds
app/__init__.py	app factory + schema bootstrap
app/models/	ORM models (one file per table)
app/routes/	Flask blueprints (HTTP layer)
chat_routes.py	/api/chat/* – message, SSE steps, sessions, feedback
database_routes.py	/api/databases/* – DB & Excel connections, Process
upload_routes.py	/api/upload/*, /api/files/* – document upload/view/delete
datasource_routes.py	/api/datasources/* – Knowledge Base list, description, Details,
document Process	
api_routes.py	/api_connectors/* – API tool save/test/delete/Process
resource_mapping_routes.py	/api/resource-mapping/* – admin sharing grants (+ LLM budgets)
auth_routes.py	/api/auth/* – register, verify, login, approve/reject
platform_routes.py	/api/platform/companies – superadmin company provisioning
model_config_routes.py	/api/model-config/* – configured LLMs, global selection
user_model_pipeline_routes	/api/user/model-pipeline, /api/models/* – per-step model choices
settings_routes.py	/api/settings/* – rag_config.yaml, query instructions
query_log_routes.py	/api/query-logs/ – Queries page
llm_call_log_routes.py	/api/llm-calls/ – LLM Calls page
export_routes.py	/api/export/results/<id> – full-result Excel download
warehouse_routes.py	/api/warehouse/* – ETL script generator
bi_semantics_routes.py	/api/bi-semantics/* – default measure per table
api_v1_routes.py	/api/v1/chat – programmatic chat endpoint
page_routes.py	HTML pages
(analytics/history/user/workspace/export-* stubs)	
app/services/	
router_service.py	★ query orchestrator (RouterService.get_smart_response)
automated_metamind.py	★ MetaMind: introspection + per-user router config
databridge_services/	★ DB track: LangGraph graph + 8 agents
spreadsheet_service.py	Parquet storage + JSON manifest
spreadsheet_query_service.py	SPREADSHEET track
spreadsheet_header_detection	finds the real header row in uploaded sheets
llm_service.py	FILES track (RAG) + document ingestion (embeddings)
api_services.py	API track (tool-calling + HTTP execution)
general_service.py	GENERAL track
data_source_finaliser.py	resolves business terms via lookup spreadsheets (pre-DB track)
column_intelligence.py	numeric-text normalisation, column role/unit, LLM column
meanings	
result_combiner.py	joins multi-source tables by value-verified keys + plan-driven
arithmetic	
bi_semantics_service.py	default measure/aggregation hints for SQL generation
model_selection_service.py	which model runs each pipeline step
model_config_access_service	LLM grants + daily budgets
llm_call_logger.py	tracked_invoke → llm_call_logs
stream_manager.py	in-memory pub/sub for live "Chain of Thought" steps
result_export_service.py	saves full result sets for Excel download
audit_service.py	log_event → audit_logs
warehouse_generator.py	warehouse ETL generation/execution
rag_config.yaml / .py	RAG + self-learning + LLM-logging settings
app/utils/	crypto (Fernet), query/llm call codes, redaction, network_guard
(SSRF), decorators	
app/templates/	Jinja HTML + vanilla JS front-end (index.html = chat)

4. Persistence model

4.1 Three logical application databases

`config/config.py` derives three database URLs from one `DATABASE_URL`. Each model chooses a database with `__bind_key__`.



"Soft" means the link is a plain integer or string column with no database foreign key, because the two tables live in different databases. The `resource_mapping`, `companies` and `users` links are real foreign keys inside `core`.

core database — identity, tenancy, access, audit

Table (model)	Key columns	Purpose
<code>companies</code> (Company)	<code>company_code</code> PK, <code>company_name</code> , <code>initial_admin_email</code>	Tenant boundary. Pre-provisioned by a superadmin.
<code>users</code> (User)	<code>id</code> , <code>email</code> (unique), <code>password_hash</code> , <code>company_code</code> FK (NULL = individual account), <code>role</code> admin/user, <code>status</code> active/pending/rejected, <code>email_verified</code> , <code>verification_token(_expires)</code> , <code>query_instructions</code> , <code>last_login</code>	Auth + tenancy. <code>query_instructions</code> = standing per-user rules injected into every query.
<code>resource_mapping</code> (ResourceMapping)	<code>company_code</code> FK, <code>resource_type</code> (database / file / api / llm), <code>resource_id</code> , <code>user_id</code> FK, <code>granted_by_user_id</code> FK, <code>daily_budget</code> , <code>budget_currency</code> ; unique (<code>resource_type</code> , <code>resource_id</code> , <code>user_id</code>)	Explicit admin-granted sharing. Nothing is shared without a row here.
<code>audit_logs</code> (AuditLog)	<code>company_code</code> , <code>user_id</code> , <code>action</code> , <code>resource_type</code> , <code>resource_id</code> , <code>details</code> JSON, <code>created_at</code>	Security/tenant event trail.

resources database — data source registry & MetaMind metadata

Table (model)	Key columns	Purpose
database_connections (DatabaseConnection)	name, type (PostgreSQL / MySQL /.../ Excel), host/port/database/username/password, config JSON, status (connected / processed / error /...), error_message, description (user), metamind_summary (AI), schema_metadata JSON (cached introspection, and for Excel rows a mirror of the manifest columns), company_code , created_by_user_id , last_tested ; config also holds column_semantics (LLM column meanings per PG table) and warehouse mapping	Every structured data source. Excel/CSV uploads are rows here too (type='Excel' , one row per sheet), but their data lives in Parquet.
files (FileResource)	document_code (unique, e.g. DOC-POL-20260101-101500), file_name, file_type, file_size, file_path, description (user), status (uploaded → processing → processed / error), error_message, metamind_summary (AI topic summary), company_code , created_by_user_id	Unstructured documents. Chunks live in Qdrant tagged with metadata.document_code .
api_connectors (ApiConnector)	integration_name (unique), base_url, endpoint, method, auth_type, api_token (Fernet-encrypted), api_description (required), status (Active), metamind_summary (redacted router-visible text), company_code , created_by_user_id	Registered REST tools for the API track.

workspace database — chat, models, learning, logs

Table (model)	Key columns	Purpose
chat_sessions (ChatSession)	session_id (unique, client-generated), title, chat_history (rendered HTML of the conversation), user_id, company_code	Sidebar chat history (one row per conversation).
query_logs (QueryLog)	query_code (QUERY00001...), user_id, company_code, question, answer, router_decision (DB / FILES / API / SPREADSHEET / GENERAL / MULTI), strategy, sources JSON, main_query (SQL, plan JSON, or METHOD URL), execution_type (fresh / reused), matched_query_code, match_score, feedback_type , remarks, related_queries JSON	One row per answered question. It is the source for the Queries page and for self-learning reuse.
response_feedback (ResponseFeedback)	user_id, company_code, query_code, question, answer, sql_query, router_decision, feedback_type like/dislike, remarks, metamind_snapshot JSON	Self-learning signal. Liked rows feed the router's feedback context.
model_configurations (ModelConfiguration)	name, model (e.g. gpt-4o-mini, api://claude-..., ollama://llama3), provider, settings JSON (custom_key, base_url, company_code, is_global_default, step_overrides ...), user_id, company_code	"Configured Models" under AI & Models.
user_model_pipelines (UserModelPipeline)	user_id, main_model, model_type_preference (oss/api), step_models JSON ({query_sense: ..., sql_generator: ...})	Per-user choice of model for each DB-pipeline step.

Table (model)	Key columns	Purpose
bi_semantics_configs (BiSemantics Config)	user_id, table_name, entity_label, measure_column, aggregation; unique (user_id, table_name)	Per-user default measure (e.g. sales orders → SUM(net_value)), layered over built-in defaults.
llm_call_logs (LLMCallLog)	call_code (LLMC00000001), user_id, company_code, session_id, query_code, purpose (router.decision, sql_generator.generate_sql, rag.answer...), provider, model, model_configuration_id, token counts, cost, duration_ms, status, prompt_preview, response_preview	Every LLM call made anywhere in the app (when llm_logging.enabled).

4.2 Non-database stores (equally important)

Store	Location	Written by	Read by
Uploaded documents	./uploads/<file>	POST /api/upload/unstructured	Process (embedding), view/download
Spreadsheet Parquet files	./uploads/spreadsheets/<connection_id>/<table>.parquet	POST /api/databases/excel, PUT /<id>/excel	SPREADSHEET track, lookup hints, finaliser, preview/download
Spreadsheet manifest	./uploads/spreadsheets/spreadsheet_metadata.json. Per table: columns with type/stats, original header label, unit, role, meaning; plus title, notes, description	spreadsheet_service.save_table, set_table_semantics (Process), delete_connection_tables, scripts/backfill_spreadsheet_metadata.py	introspect_spreadsheets, SPREADSHEET track
Qdrant collection saarthi_unstructured	Qdrant service	Document Process (llm_service.process_to_embeddings)	FILES track, introspect_qdrant
External DB table comments	COMMENT ON TABLE in the customer's Postgres	PostgreSQL Process (enrich_table_descriptions_with_llm)	introspect_databridge_db (table descriptions)
Result exports	<instance>/result_exports/<uuid>.json (7-day TTL)	chat_routes.send_message → save_result_export	GET /api/export/results/<id> (Excel download)
rag_config.yaml	app/services/rag_config.yaml	Settings page (POST /api/settings/rag-config)	RAG, self-learning switch, LLM-logging pricing

Store	Location	Written by	Read by
Live step stream	In-process memory (<code>stream_manager</code>)	Every track, while a query runs	GET <code>/api/chat/stream_steps</code> (SSE)
Process logs	<code>./logs/process_<id>_<ts>.log</code>	<code>run-agentic-process</code>	Operators only
Browser <code>localStorage</code>	<code>saarthi_chat_history</code> , JWT	Front-end	Front-end

5. Tenancy & visibility rule ("own + granted")

Every resource-listing and MetaMind function applies the same rule:

```
visible(user, resource_type) =
  rows WHERE created_by_user_id = user.id
  U  rows WHERE id IN (SELECT resource_id FROM resource_mapping
                       WHERE resource_type = :type AND user_id = user.id)
```


- Sharing a `company_code` **does not** make a resource visible. Only an admin's `resource_mapping` grant does.
- Implemented in `automated_metamind._visible_resource_ids()` (returns the granted IDs; callers OR them with "own"), `visible_document_codes()`, `_visible_postgresql_connections()`, `datasource_routes.get_visible_datasources()` and the list endpoints in each route file.
- **Modify/delete** rights are wider: the creator, **or** an admin of the resource's company (`_can_modify_connection`, file/API delete checks).
- Chat sessions and query-log visibility are per user. The query log is per company when the user has a `company_code` (`query_log_routes`).

6. MetaMind — the per-user router config

"MetaMind" is the description of *what this user can query*. It is **not stored as one document**.

`automated_metamind.generate_router_config(user_id, sap_db_config=None, use_cached_metadata=False)` recomputes it every time, from the tables listed above. (A legacy `router_configs` table used to cache it and is now dropped at boot.)

6.1 What `generate_router_config` does (in order)



6.2 What `introspect_databridge_db(db_config)` captures per table

For each `public` base table, skipping Saarthi's own tables and any `_`-prefixed padding tables:

- `description`: `obj_description()` (the table COMMENT), or the fallback "Table storing X records."
- `row_count`: `COUNT(*)`
- `constraints`: primary key, foreign keys, unique columns
- per column: `data_type`, `nullable`, and, when `row_count ≤ 100,000`, also `unique_values`, `null_count` and up to 5 `sample_values`
- then `_apply_db_column_intelligence` adds a deterministic `role` (identifier/measure/dimension/date/flag/text) and `unit` per column (`column_intelligence.profile_column`), and merges any LLM column `meaning` stored in `connection.config.column_semantics`. A stored meaning is only applied if the table's column-name signature still matches.

This is the shape stored in `DatabaseConnection.schema_metadata`. The connection's free-text `description` is prefixed onto each table description at read time, so editing the description takes effect without re-introspecting.

6.3 Two modes and who calls which


Mode	Cost	Side effects	Callers
Live (<code>use_cached_metadata=F</code> <code>alse</code> , default)	Connects to every visible PG connection: <code>COUNT(*)</code> , then per-column profiling	Updates <code>database_connections.status/error_message/schema_metadata/metamind_summary</code> , <code>api_connectors.metamind_summary</code> , Excel <code>metamind_summary</code>	DB create/update/ Process , Excel upload/edit/ Process , document Process , API save/ Process , Knowledge Base Details popup (<code>/metadata</code>)
Cached (<code>use_cached_metadata=T</code> <code>rue</code>)	Reads stored JSON only (Qdrant/API/spreadsheets parts are still read live)	None for DB. API/spreadsheet <code>metamind_summary</code> writes are no-ops when the text is unchanged	Every chat question (<code>router_service._load_route_r_config</code>), DB-track schema fallback, feedback snapshot

Key consequence: a PostgreSQL connection contributes **no tables** to routing until it has been introspected live at least once. Creating it or clicking **Process** does that. Schema changes in the customer database are invisible to Saarthi until the next live introspection.

7. Setup flows (what populates the tables)

Every flow below lists the **module call sequence** and the **table/store writes**. In the sequence diagrams,  means "writes".


7.1 Company provisioning, signup, login



Login is refused when `email_verified` is false or `status` is `pending` / `rejected`. The JWT identity is `user.id`. `get_current_user()` resolves it on every protected route.

7.2 PostgreSQL connection — Save → Process

UI: **Database Connections** page (`database_connections.html`).



What Process changes for later queries:

- The cached `schema_metadata` (tables, columns, samples, constraints) is now fresh. **This is the schema the router and the SQL agents see.**
- Table descriptions improve, because LLM-written COMMENTS replace the generic "Table storing X records." The router uses these descriptions to pick tables.

- Every column gets a role/unit and an LLM-written meaning. The router, SQL agents and result combiner all see these.
- `metamind_summary` is refreshed. It appears in the Knowledge Base list and Details popup.


Edit (`PUT /api/databases/<id>`) re-runs the live config for owner, editor and grantees. **Delete** removes the `resource_mapping` rows and the connection row. Nothing else needs cleaning up, because routing is computed live and the connection simply disappears from the next query.

Advanced: `POST /api/databases/run-agentic-process/<id>` runs two scripts as subprocesses: `databrige_services/db.py` (seeds an SAP-style demo dataset into the target DB) and `metamind.py` (the file no longer exists, see §16). It then introspects, stores `schema_metadata`, regenerates config for everyone affected, and writes `audit_logs (agentic_process_run)`. The UI no longer calls it.

Note on non-PostgreSQL types (MySQL, Oracle, ...). Rows can be saved, but `_visible_postgresql_connections` only introspects `type == 'PostgreSQL'`. These connections never contribute schema, so the DB track cannot query them.

7.3 Excel / CSV upload → Process

UI: **Spreadsheets** page (`spreadsheets.html`) or the Database Connections page (Excel type).




What Process changes for later queries: the manifest `description` (and per-column meanings) fills `SPREADSHEET.tables[t].description` in the routing menu. That description is the main signal the router uses to choose `query_spreadsheet_data`, and it also feeds the **lookup-hint** cross-reference (§9.1) that lets DB questions use spreadsheet code tables.

Edit (PUT `/api/databases/<id>/excel`) renames, changes the description, and/or replaces the data with a single-sheet file. On a data replace it overwrites the Parquet file and manifest entry, **sets** `metamind_summary=NULL` and `status='connected'` (so Process must be run again), then regenerates config. **Delete** removes the mapping rows, the connection row, the Parquet files and the manifest entries. **Preview/Download** read the Parquet data.

7.4 Document upload → Process (FILES)

UI: **Unstructured Data** page (`unstructured_data.html`).



What Process changes for later queries: the document's chunks become retrievable, and `metamind_summary` appears in `FILES.vector_store_info.documents[].description`. The router reads this to decide on `search_documents` and to pick `document_codes`.

Delete (DELETE `/api/files/<code>` or `/api/datasources/unstructured/<code>`): purge the Qdrant points by `document_code` → remove the disk file → delete the `resource_mapping` rows → delete the `files` row → write `audit_logs`.

7.5 API connector Save → Process (API)

UI: **REST APIs** page (`api_connectors/rest_apis.html`).

```

POST /api_connectors/test_connection → network_guard.is_safe_url (SSRF block) → live HTTP ping,
no writes
POST /api_connectors/save_tool → is_safe_url → upsert api_connectors (token Fernet-
encrypted;
                                empty token on edit keeps the old one) → audit_logs
                                → generate_router_config(user) [LIVE] → 🗑️
api_connectors.metamind_summary
POST /api_connectors/tools/<name>/process → generate_router_config(uid) for owner, clicker,
grantees
                                (refreshes metamind_summary; no other state)
DELETE /api_connectors/delete_tool/<name> → resource_mapping rows + api_connectors row →
audit_logs

```

At query time the API track reads the connectors through `api_services.fetch_and_translate_tools()`. That function turns each **Active** connector into an OpenAI function-tool schema (name = sanitised `integration_name`, description = redacted `api_description`).

7.6 Knowledge Base page — description edit & Details

- GET `/api/datasources/` → `get_visible_datasources(user)`: a merged list of DB, Excel, file and API resources showing `status`, `description`, `metamind_summary`, `creator`, and `owner/shared`.
- PUT `/api/datasources/<type>/<id>/description` → writes `description` (or `api_description`). No regeneration is needed, because the next query reads it live (DB descriptions are prefixed at read time).
- GET `/api/datasources/<type>/<id>/metadata` (**Details** popup) → `generate_router_config(user)` **in live mode**, sliced down to the selected resource, returning both summary and full metadata. Note that this re-scans the user's PG connections.

7.7 Resource Mapping (sharing) & LLM budgets

UI: **Resource Mapping / LLM Mapping** pages. Admin-only (`@admin_required`), and limited to users and resources of the admin's own company.

```

POST /api/resource-mapping → single grant → 🗑️ resource_mapping →
audit_logs(resource_mapping_grant)
POST /api/resource-mapping/bulk → users × resources → 🗑️ resource_mapping (skips duplicates) →
audit_logs
PATCH /api/resource-mapping/<id>/budget → 🗑️ daily_budget / budget_currency (llm grants only)
DELETE /api/resource-mapping/<id> → 🗑️ delete → audit_logs(resource_mapping_revoke)

```


No regeneration happens. The grantee's next question resolves visibility live. Because `schema_metadata` is stored per connection, a grantee immediately sees the cached schema of a shared DB connection that has already been processed.

7.8 Model configuration & pipeline selection

- `model_config_routes` → CRUD on `model_configurations`. It seeds default open-source models per company and stores `custom_key` / `base_url` / `is_global_default` / `step_overrides` in `settings`.
- `user_model_pipeline_routes` → POST `/api/user/model-pipeline` upserts `user_model_pipelines` (`main_model`, `step_models`). `/auto-fill` and `/models/recommendations` produce presets from `model_registry_service`.

- `settings_routes` → `POST /api/settings/query-instructions` writes `users.query_instructions`.
`POST /api/settings/rag-config` rewrites `rag_config.yaml` (RAG knobs, `self_learning.enabled`, LLM pricing).
- `bi_semantics_routes` → upsert/delete `bi_semantics_configs`.

7.9 Status lifecycles



Re-processing a document calls `from_documents` again **without deleting the old points first**, so its chunks are duplicated in Qdrant (see §16).


8. Query flow — what happens when a question is asked

8.1 Front-end handshake

`index.html` does the following, in order:

1. Opens an `EventSource` on `GET /api/chat/stream_steps?session_id=...&token=<JWT>`. This route accepts the JWT from the query string only.
2. Once it receives `onopen` (the server immediately sends `{connected:true}`), it sends `POST /api/chat/message` with `{message, session_id, model_name, custom_key?, system_instructions?}`.
3. It renders live step cards from the SSE events until it sees `step == "DONE"`.
4. It renders the POST JSON (answer, table, chart, insights, steps, `query_code`, `export`).
5. It saves the conversation HTML through `POST /api/chat/sessions` → `chat_sessions` (and to `localStorage`).
6. Thumbs up/down → `POST /api/feedback` → \$10.

8.2 End-to-end sequence



8.3 Layer-by-layer reference (`router_service.py`)

Layer	Function(s)	Reads	Writes	LLM calls
L1 heuristic	<code>classify_query_heuristic</code>	<code>general_knowledge_config.js</code> on	<code>query_logs</code> (on hit)	1 (general answer)
L1.5 reuse	<code>_find_reusable_query</code> → <code>REUSE_DISPATCH[DB SPREADSHEET]</code>	<code>query_logs</code> (<code>feedback_type='like'</code> , <code>router_decision ∈ {DB,SPREADSHEET}</code> , <code>main_query</code> not null, scoped to company or user, newest 200); HF MiniLM embeddings	<code>query_logs</code> (reused)	0 for routing/SQL. DB reuse still runs the insight generator in the formatter
L2 context	<code>_load_router_config</code> → <code>generate_router_config(cached)</code> ; <code>fetch_and_translate_tools</code> ; <code>_build_feedback_context</code> ; <code>_build_router_messages</code>	<code>database_connections.schema_metadata/description</code> , <code>resource_mapping</code> , <code>files</code> , Qdrant scroll, manifest, Parquet (lookup hints), <code>api_connectors</code> , <code>response_feedback</code>	–	0
L3 routing	<code>ChatOpenAI("gpt-4o-mini").bind_tools(_ALL_TOOLS)</code> via <code>tracked_invoke</code>	–	<code>llm_call_logs</code>	1
L4 dispatch	<code>_execute_tool_call(s)</code> / <code>TOOL_DISPATCH</code> , using a <code>ThreadPoolExecutor</code> with its own Flask app context when >1	see §9	see §9	per track
L5 combiner + synthesis	<code>result_combiner.combine_results</code> (or <code>_pick_primary_tabular_result</code>), <code>_decide_output_format</code> , <code>_generate_chart_for_merged_table</code> , <code>_build_multi_strategy</code>	<code>answer_guidelines.md</code>	<code>query_logs</code> (MULTI), <code>llm_call_logs</code>	1 combiner plan + 1 synthesis (+1 chart if needed)

Router tools (the LLM chooses one or more, and each call carries a rewritten sub-question):

Tool	Args	Dispatches to	<code>router_decision</code> logged
<code>check_data_source_status</code>	<code>track</code>	<code>_answer_status_check</code> (reads the config only)	<i>not logged</i>
<code>query_database</code>	<code>question</code> , <code>tables[]</code>	<code>_run_db_track</code>	DB
<code>query_spreadsheet_data</code>	<code>question</code> , <code>tables[]</code>	<code>_run_spreadsheet_track</code>	SPREADSHEET
<code>search_documents</code>	<code>question</code> , <code>document_codes[]</code>	<code>_run_files_track</code>	FILES
<code>call_external_api</code>	<code>question</code> , <code>tool_name</code>	<code>_run_api_track</code>	API
<code>answer_general_knowledge_tool</code>	<code>question</code>	<code>_run_general_track</code>	GENERAL
(no tool)	–	the router's own text	GENERAL

Outcomes: 0 tools → direct answer. 1 tool → that track's result returned **unchanged**. ≥2 tools → merged tables + synthesised prose. A failed track is excluded from synthesis and named in a trailing "Note: I couldn't get data from ...".

8.4 Response contract (`response` object)

```
{
  "answer": "...", "sql": "SELECT ..." , "table": [ {...} ], "chart": {"bar": {...}, "line": {...}, "pie": {...}, "recommended": "bar"},
  "insights": [...], "format": "text|kpi|table|chart", "steps": ["Title - description", ...],
  "chain_of_thought": [...], "router_decision": "DB|FILES|API|SPREADSHEET|GENERAL|MULTI",
  "execution_type": "fresh|reused", "query_code": "QUERY00042",
  "strategy": "...", "sources": ["orders<Database>", ...], "main_query": "...",
  "export": {"export_id": "...", "row_count": N, ...}
}
```


9. Track internals

9.1 DB track — `_run_db_track` → LangGraph

Pre-step: Data Source Finaliser (`data_source_finaliser.finalize_data_source_strategy`). For every DB column carrying a `lookup_hint` (added by MetaMind when the column's sample values are all codes in a small spreadsheet column), it tries to resolve candidate terms from the question ("copper", "Europe") against that lookup table's label columns (reading the **Parquet** data). Matches are prepended to the question as a "KNOWN CODE TRANSLATIONS" block, e.g. `marc.material_group: "copper" -> 'RM03'`.

`run_data_bridge_agent(...)` (`databridge_services/langgraph_agent.py`):

1. `schema = _build_schema_for_user(user_id, router_config) → to_sql_agent_schema(DB.tables, inferred_relations)`. This reuses the router config already computed in L2, so **the schema is the cached schema_metadata**.
2. `db_config = resolve_query_execution_config(user_id) → the first visible PostgreSQL connection with its password decrypted`. If there is none, it falls back to the env `DATABRIDGE_DB_*` / app DB.
3. Builds request-scoped `QuerySimplifierAgent`, `QuerySenseAgent` and `QueryValidatorAgent` with that schema. `SQLGenerator`, `QueryFormatter`, `DataInsightGenerator`, `DataVisualizer` and `ErrorDiagnosis` are module singletons.
4. Streams `langgraph_app.stream(initial_state)`. After each node it pushes a human-readable step card, and at the end pushes `"DONE"`.



Node	Reads (DB/stores)	LLM purpose
model choice for every node	user_model_pipelines, model_configurations via get_model_for_step	-
simplifier	schema	query_simplifier.simplify
query_sense	schema, hint_tables, feedback	query_sense.plan

Node	Reads (DB/stores)	LLM purpose
validator	schema (column/alias checks)	–
sql_generator	bi_semantics_configs + built-in defaults (build_measure_guidance)	sql_generator.generate_sql , sql_generator.narration
query_formatter	customer Postgres (read-only SELECT)	data_insight.generate
insight_generator	result rows	data_insight.generate
visualizer	result rows; top-N from query_instructions	–

The track returns `main_query = SQL`, which `query_logs.main_query` stores. That stored SQL is what enables later **reuse** (§10).

9.2 SPREADSHEET track — `spreadsheet_query_service.answer_from_spreadsheets`

```

available_tables = spreadsheet_service.list_all_tables()           # whole manifest (see §16)
narrow to hint_tables if given
if question is about a table's metadata (rows/columns/upload date) → answer from manifest, no LLM
plan
plan = LLM(purpose spreadsheet.plan): JSON {tables, filters, group_by, aggregations, join, sort,
limit}
plan = _validate_plan(plan, available_tables)                     # unknown table/column →
PlanValidationError
df    = _execute_plan(plan)                                       # pandas over Parquet files
answer= LLM(purpose spreadsheet.answer): analyst-style summary of df sample
return {answer, table: df records, tables, plan}

```

`_run_spreadsheet_track` stores `main_query = json.dumps(plan)`. A liked spreadsheet answer can then be replayed by `_run_reused_spreadsheet_query` (it re-runs `_execute_plan` and returns the old answer text verbatim).

9.3 FILES track — `llm_service.answer_from_docs`

```

doc_codes = visible_document_codes(user)                         # files own + granted
if none → "You don't have any documents..."
narrow to router's document_codes hint if all are visible
Qdrant filter: metadata.document_code ∈ doc_codes
LLM intent-analysis line (rag.intent_analysis) – shown as a step card
queries = [q] (+ multi_query variations, + HyDE doc – off by default)
similarity_search top_k=3 per query → merge/dedupe
LLM answer strictly from context (rag.answer) – gpt-4o-mini / gpt-4o / llama3 / api:// /
ollama://
return {answer, document_codes, rag_chain_of_thought}

```

The track maps `document_codes` back to `files.file_name` for `sources` and `strategy`. Nothing re-runnable is stored in `main_query`.

9.4 API track — `api_services.ask_dynamic_model_with_tools`

```
tools = fetch_and_translate_tools()          # Active api_connectors → function schemas
LLM.bind_tools(tools) (forces hint tool_name when supported)  # OpenAI / Ollama /api/chat /
api://
if a tool is chosen:
    connector = api_connectors by integration_name
    headers/params from auth_type + decrypt(api_token) (Bearer, or X-API-Key header + api_key
param)
    network_guard → HTTP request → response text
    LLM summarises response
return {answer, tool_call:{tool_name, method, url}, steps}
```

`query_logs.main_query` = "METHOD URL". The API track is never reused.

9.5 GENERAL track & status check

- `general_service.answer_general_knowledge`: a system prompt with today's date/time, the user's instructions and feedback context → the chosen model. Logged as `GENERAL`.
- `_answer_status_check`: reads the routing menu only and describes which tracks have data (e.g. "2 database tables, 3 documents..."). No agents run, and nothing is written to `query_logs`.

9.6 Multi-source combine (`result_combiner.combine_results`)

Runs when at least two successful tracks returned rows. Example: "total disbursed amount, and the commission at the lowest slab rate" needs DB rows × spreadsheet rates.


```
for each track table → _clean_frame: numbers-as-text → numbers + units, drop stray note rows;
alias db/sheet/api/docs
find_key_candidates: join keys verified by VALUE overlap (≥ 30%), not just same column names
LLM (router.result_combiner) proposes a small JSON plan: join key pair + how (left/inner/outer),
row-wise min/max/mean/sum, arithmetic (multiply/divide/add/subtract/percent_of)
validate_plan against the real columns and verified keys → on failure: _fallback_plan (best
key, no derived columns)
execute_plan with fixed pandas code (the LLM never writes code) → combined table + grand totals +
join_stats
```

The result feeds the synthesis prompt as "[Combined result - computed exactly by code ...]", so the LLM reports figures instead of computing them. With no combination (fewer than 2 tables, or the combiner fails) the router falls back to `_pick_primary_tabular_result`. When the sources share no matching values, the combiner's notes become a "Data alignment note" so the LLM does not substitute one source's field for another. The output format is decided by row and column counts. If it is `chart` and no track supplied a chart, one is generated from the combined table.

10. Self-learning (feedback & query reuse)

Enabled by `rag_config.yaml` → `self_learning.enabled: true`.

10.1 Writing feedback



- The Settings → Self-Learning page edits feedback afterwards. `PUT /api/chat/feedback/<query_code>` upserts the `response_feedback` row and mirrors it to `query_logs`. `DELETE` clears both.
- Scope: feedback is **company-wide** when the user has a `company_code`, otherwise per user.

10.2 Reading feedback at query time

Mechanism	Source table & filter	Threshold	Effect
Reuse (L1.5)	<code>query_logs</code> with <code>feedback_type='like'</code> , <code>router_decision ∈ {DB, SPREADSHEET}</code> , <code>main_query</code> not null	<code>max(0.94, threshold from instructions)</code> cosine on MiniLM embeddings	Skips the router LLM and SQL/plan generation entirely. Re-executes the stored SQL/plan live. Logged as <code>execution_type='reused'</code> , <code>matched_query_code</code> , <code>match_score</code>
Feedback context (L2)	<code>response_feedback</code> with <code>feedback_type='like'</code> (newest 200)	default 0.80 , overridable via Query Instructions such as "Query match threshold 90%" (<code>instruction_settings.match_threshold_from_instructions</code>)	Adds "a similar question was LIKED before" lines to the router prompt and to each track's prompts. The matches are stored in <code>query_logs.related_queries</code>

Disliked rows are stored but **never** matched. This is deliberate, to avoid injecting unrelated "avoid this" remarks into other questions.

11. Observability: Query log, LLM call log, Audit log

Log	Written by	Granularity	UI
query_logs	router_service._log_query (every path except status checks and crashes)	1 row per answered question	Queries page /query_logs → GET /api/query-logs/?source=&feedback=&user_id= (company-wide for company users)
llm_call_logs	llm_call_logger.tracked_invoke (LangChain) / track_ollama_call / record_ollama_call (raw Ollama)	1 row per LLM call, with tokens, cost from rag_config.llm_logging.pricing, duration, prompt/response previews	LLM Calls page /llm_calls → GET /api/llm-calls/, /summary (grouped by purpose)
audit_logs	audit_service.log_event	Security/tenant events: register, login(_blocked), approve/reject, resource created/updated/deleted/processed, mapping grant/revoke, agentic run	Audit Logs page

Common purpose values: router.decision, router.result_combiner, router.multi_source_synthesis, query_simplifier.simplify, query_sense.plan, sql_generator.generate_sql, sql_generator.narration, data_insight.generate, spreadsheet.plan, spreadsheet.answer, rag.intent_analysis, rag.answer, rag.hyde, rag.multi_query, rag.document_summary, rag.image_caption, metamind.table_description, metamind.column_semantics.

12. Model selection & LLM backends

- **Router, synthesis, and the MetaMind/Process summaries are always `gpt-4o-mini`** (OpenAI key from `custom_key` or `OPENAI_API_KEY`), whatever model the user picked.
- **Track models** follow the chat dropdown `model_name`:
 - `gpt-4o`, `gpt-4o-mini` → LangChain `ChatOpenAI`
 - `llama3` / `ollama://<model>` → raw HTTP to `http://ollama:11434/api/generate` (DB agents) or `/api/chat` (API track)
 - `api://<model>` → OpenAI-compatible endpoint, where `custom_key` / `base_url` come from the matching `model_configurations.settings`
- **Per-step override for the DB pipeline** (`model_selection_service.get_model_for_step`):
 1. the chat-bar model, if it differs from the user's saved `main_model` (explicit per-request override)
 2. `user_model_pipelines.step_models[step]`
 3. `user_model_pipelines.main_model`
 4. `step_overrides[step]` on the global default `model_configurations` row
 5. the global default model, or the requested model
- **Budgets:** `resource_mapping.daily_budget` for `llm` grants, checked by `model_config_access_service.get_budget_status`. Spend is summed from `llm_call_logs.cost` by `model_configuration_id` (see §16 on why it currently stays 0).

13. Data warehouse builder (side feature)

`/build_warehouse` and `/warehouse/mapping` pages + `warehouse_routes.py` + `warehouse_generator.py`:

- **Sources:** DB-track tables (via MetaMind) and SPREADSHEET tables (Parquet).
- **Target:** a `database_connections` row with `config.role == "warehouse_target"`.
- **Mapping state** is stored inside that target row's `config` JSON: `warehouse_table_groups` (1 source = 1:1 copy, several = INNER/LEFT join) and `warehouse_mapping[target_table]` (rename/retype/exclude/transform per column).
- `POST /api/warehouse/generate` emits a standalone Python ETL script. `POST /api/warehouse/run` executes it in-process. Run history is written to a `warehouse_sync_log` table inside the **target** database. Health checks and a data-model view are also available.

14. Table x flow CRUD matrix

C = create, R = read, U = update, D = delete. "Manifest" and "Qdrant" are the non-database stores.

Flow ↓ / Table →	companies	users	resource_mapping	audit_logs	database_connection	files	api_connectors	Manifest/Parquet	Qdrant	chat_sessions	query_logs	response_feedback	model_configurations	user_model_pipeline	bi_semantics_configs	llm_call_logs
Provision company	C															
Register / verify / login / approve	R	C U		C												
DB connection save/edit		R	R	C	C U (schema_metadata, summary, status)											
DB Process		R	R		U (status, schema_metadata, summary, config.column_semantics) + COMMENT ON TABLE on external DB											C
Excel upload / edit		R	R	C	C U			C U								
Excel Process		R	R		U (status, summary, schema_metadata)			U (description, column meanings)								C
Document upload		R		C		C										
Document Process		R	R	C		U (status, summary)			C							C
API save / Process		R	R	C			C U (summary)									
Delete DB/Excel/doc/API		R	D	C	D	D	D	D	D							
Resource mapping grant/revoke/budget		R	C U D	C	R	R	R						R			
Description edit / Details		R	R		U / R(+live introspect)	U / R	U / R	R	R							
Ask question (fresh)		R	R		R (schema_metadata, desc)	R	R	R	R		C	R	R	R	R	C
Ask question (reused)		R			R			R			C + R		R			C (insights)
Save chat session		R								C U						
Like / Dislike / edit feedback		R	R		R	R	R	R	R		U	C U D				
Model config / pipeline / BI semantics		R											C U D	C U	C U D	
Queries / LLM Calls pages		R									R					R

15. Configuration reference

Setting	Where	Effect
DATABASE_URL	env	Host/credentials for the 3 app databases (saarathi_core_db , saarathi_resources_db , saarathi_workspace_db)
SECRET_KEY , JWT_SECRET_KEY	env	Required in production (startup refuses the defaults)
ENCRYPTION_KEY	env	Fernet key for API tokens (and any value passed through crypto.encrypt)
OPENAI_API_KEY	env	Router, synthesis, summaries, OpenAI-backed tracks
REDIS_URL	env	Rate-limit storage
ALLOWED_ORIGINS	env	CORS allow-list (empty → cross-origin denied)
SUPERADMIN_EMAILS	env	Who may provision companies
QDRANT_URL , QDRANT_COLLECTION	env / rag_config.y aml	Vector store
DATABRIDGE_TARGET_*	env	Legacy dev fallback DB for MetaMind when the user has no PG connection
DATABRIDGE_DB_* / PG*	env	Fallback SQL execution DB for QueryFormatterAgent
rag_config.yaml	file / Settings page	embedding model, chunk size/overlap, top_k, HyDE, multi-query, table/image extraction, self_learning.enabled , llm_logging (+ pricing)
general_knowledge_config.json	file	L1 heuristic patterns
answer_guidelines.md	file	Style rules injected into synthesis
users.query_instructions	DB / Settings	Standing per-user rules: top-N for charts, match threshold, ranking defaults
Rate limits	code	200/min default; chat message 20/min; uploads 20/min; API tool save/test/process 20/min

16. Observed gaps & risks in the current code

These came up while tracing the flows above. They are recorded here so the design reflects reality. Each one is a candidate follow-up, not something this document changes.

#	Area	Observation	Impact
1	API track	<code>_run_api_track</code> calls <code>ask_dynamic_model_with_tools(..., feedback_context=..., hint_tool_name=...)</code> , but the function signature (<code>api_services.py:90</code>) accepts neither keyword. Its body also references <code>hint_tool_name</code> .	Every API-track call raises <code>TypeError</code> . <code>_execute_tool_call</code> catches it and returns an error result, so the API track effectively never answers.
2	Tenancy	<code>fetch_and_translate_tools()</code> loads all Active <code>api_connectors</code> , and <code>answer_from_spreadsheets</code> uses the whole manifest (<code>list_all_tables()</code>). Neither is scoped by "own + granted".	Router context is scoped, but these execution paths can reach other tenants' API tools and spreadsheet tables.
3	Conversation memory	<code>chat_routes.send_message</code> never passes <code>chat_history</code> to <code>get_smart_response</code> .	The router is stateless per message; follow-up questions ("and for last year?") lack context.
4	Credentials	<code>database_connections.password</code> is stored as typed . <code>create / update</code> never call <code>encrypt()</code> , although the model docstring says it is encrypted (<code>decrypt()</code> passes plaintext through).	DB passwords are stored in plaintext.
5	Multi-DB	<code>resolve_query_execution_config</code> executes SQL against the first visible PG connection, while the schema merges tables from all of them.	With 2+ PG connections, SQL for tables of the 2nd connection fails.
6	Non-PG connections	Only <code>type == 'PostgreSQL'</code> is introspected or executed.	MySQL/Oracle/... connections can be saved and "processed" but are never queryable.
7	Re-processing documents	Process re-indexes without first deleting the old points for that <code>document_code</code> .	Duplicate chunks in Qdrant, skewing retrieval and chunk counts.
8	Agentic process	<code>run-agentic-process</code> invokes <code>databridge_services/metamind.py</code> , which no longer exists.	Endpoint fails at step 2 (UI no longer calls it).
9	LLM budgets	No caller passes <code>model_configuration_id</code> (or <code>query_code</code>) to <code>tracked_invoke</code> .	Daily LLM budgets never accrue spend, and LLM calls cannot be joined to their query.
10	Identity fallback	<code>_resolve_feedback_user</code> falls back to user id 1 when the JWT cannot be resolved.	Activity could be attributed to user 1.
11	Details popup	<code>/api/datasources/<type>/<id>/metadata</code> runs a live introspection (<code>COUNT(*)</code> and profiling) for all of the user's PG connections on every open.	Slow on large schemas, and it rewrites status as a side effect.

#	Area	Observation	Impact
1 2	Stream manager	In-process memory, keyed by <code>session_id</code> .	With several gunicorn workers, the SSE request and the POST can land on different workers, and steps are lost. Needs Redis pub/sub for multi-worker setups.
1 3	Debug logging	<code>get_smart_response</code> prints the full inputs, including a preview of the <code>custom_key</code> .	Secrets and PII in stdout logs.